

Diabetes Prediction with Pima Diabetes Dataset using Decision Tree and Neural Network (From Abdul-Qadir)

```
In [2]: # importing libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
```

```
In [3]: # Random Seed for Reproducibility
np.random.seed(42)
tf.random.set_seed(42)
```

```
In [4]: # Loading Data
df = pd.read_csv('pima_diabetes_data.csv')
print('Dataset shape:', df.shape)
print(df.describe())
```

Dataset shape: (768, 9)

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
In [5]: # Preprocessing (replace 0 with NaN, fill with mean)
cols = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']
df[cols] = df[cols].replace(0, np.nan)
df.fillna(df.mean(), inplace=True)
```


```
In [6]: # Features/target, train/test split, scaling
X = df.drop('Outcome', axis=1)
y = df['Outcome']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
scaler = StandardScaler())
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Decicion Tree

```
In [7]: # Decision Tree
dt = DecisionTreeClassifier(max_depth=4, random_state=42)
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
print('\nDecision Tree Accuracy:', accuracy_score(y_test, y_pred_dt))
print(classification_report(y_test, y_pred_dt))
sns.heatmap(confusion_matrix(y_test, y_pred_dt), annot=True);
plt.title('DT Confusion Matrix');
plt.show()
```

Decision Tree Accuracy: 0.7857142857142857

	precision	recall	f1-score	support
0	0.83	0.84	0.84	100
1	0.70	0.69	0.69	54
accuracy			0.79	154
macro avg	0.76	0.76	0.76	154
weighted avg	0.78	0.79	0.79	154



Neural Network

```
In [10]: # Neural Network
model = keras.Sequential([
    layers.Dense(16, activation='relu', input_shape=(X_train.shape[1],)),
    layers.Dense(8, activation='relu'),
    layers.Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model.fit(X_train, y_train, epochs=50, batch_size=16, validation_split=0.2, verbose=0)
loss, acc = model.evaluate(X_test, y_test, verbose=0)

print('\nNeural Network Accuracy:', acc)

y_pred_d1 = (model.predict(X_test) > 0.5).astype("int32").flatten()

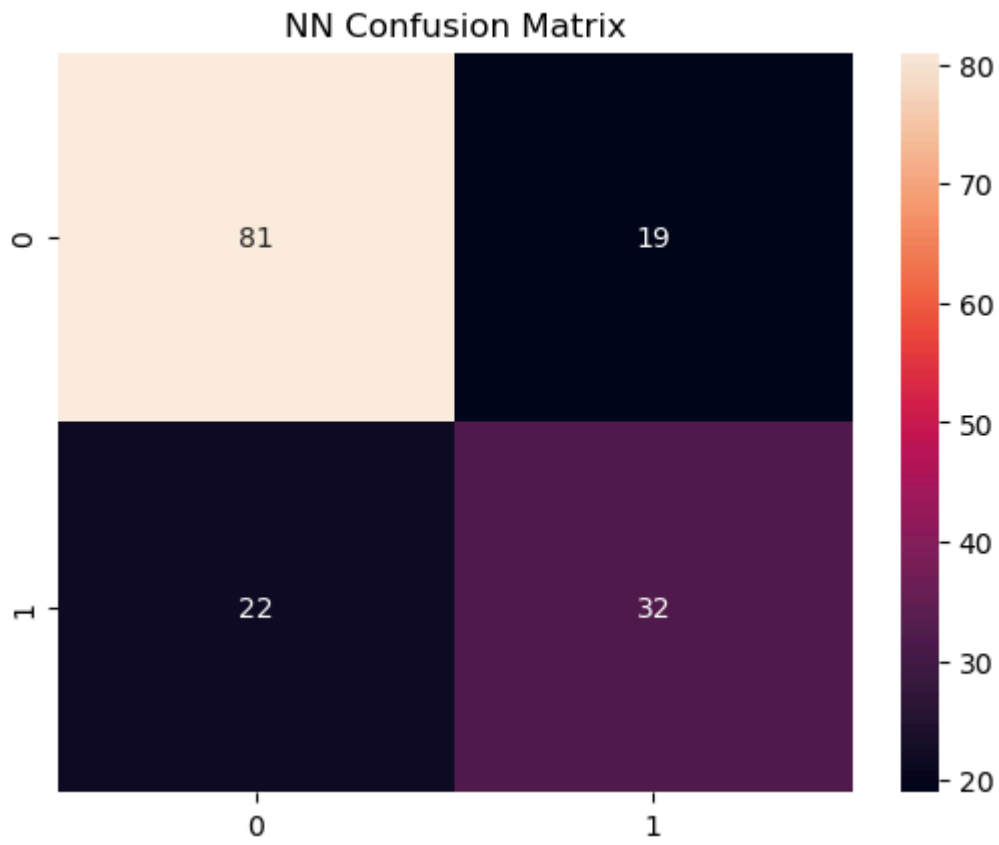
print(classification_report(y_test, y_pred_d1))

sns.heatmap(confusion_matrix(y_test, y_pred_d1), annot=True);
plt.title('NN Confusion Matrix');
plt.show()
```

Neural Network Accuracy: 0.7337662577629089

5/5 ————— 0s 37ms/step

	precision	recall	f1-score	support
0	0.79	0.81	0.80	100
1	0.63	0.59	0.61	54
accuracy			0.73	154
macro avg	0.71	0.70	0.70	154
weighted avg	0.73	0.73	0.73	154



In []: